Significant Changes from Previous Version

In response to reviewer feedback about incremental novelty and design motivation, we have made the following substantial improvements:

1. **Unified Theoretical Framework**

   We have fully developed the ***guest-delegation*** principle as a cohesive theoretical framework underlying our approach, clarifying how this represents a paradigm shift rather than an incremental improvement over existing hyperivsor-based and kernel-based solutions.
2. **Strengthened Motivation**

   We have added detailed empirical analysis in Section 2.3 demonstrating the fundamental limitations of hypervisor-based approaches, including:
   - Quantitative measurement of TLB flushes and overhead (Tbl. 1)
   - Analysis of EPT-friendly PEBS scalability issues (Fig. 2)
3. **Expanded Evaluation**

   We have addressed key evaluation gaps by:
   - Adding direct comparisons with hypervisor-based approaches (TPP-H) across seven real-world workloads
   - Including latency analysis for interactive applications (Fig. 12), showing 23% reduction in tail latency
4. **Clarified Novelty of PEBS Approach**

   We have substantiated our claims about EPT-friendly PEBS with:
   - Technical explanation of previous architectural limitations (Section 2.3.2)
   - Analysis of isolation properties ensuring multi-tenant security
   - Context-switch-based sampling implementation that reduces CPU overhead by 16×
5. **Added Threat Model Discussion**

   Despite this is not the main focus of our paper, we have included a threat model in Section 7.2, addressing security considerations for untrusted guest code running in multi-tenant environments.